

# CS 511: Homework Assignment 1

## Due: Tuesday 15 September, 11:59pm

### 1 Assignment Policies

**Collaboration Policy.** This assignment must be done individually or in groups of two. It is acceptable for students to collaborate in understanding the material but not in solving the problems or programming, unless they are members of the same group. Use of the Internet is allowed, but should not include searching for existing solutions.

**Under absolutely no circumstances code can be exchanged between students of different groups.** Excerpts of code presented in class can be used.

**Assignments from previous offerings of the course must not be re-used.** Violations will be penalized appropriately.

**Late Policy.** Late submissions are allowed. The policy is 2 points off for every hour past the deadline.

### 2 Concurrent File Writing

The idea behind this assignment is to get students familiar with the basics of threads in Java. The task itself is quite simple, but the implementation does involve some consideration about how multi-threaded programs work, and forces the student to think about the role each thread should play in the task.

In this assignment, students will implement a method that reads a text file into memory, concurrently rearranges this buffer, and then writes the rearranged buffer into an new output file. Figure 1 depicts the architecture of this system. A file is read into an input buffer. Then some number of threads (in this figure there are four of them, however this is just an example; your solution may require more or less than four threads; please read on) each read disjoint pieces of the input buffer and write them into disjoint pieces of the output buffer. The input and output buffers are different buffers, in other words, they are allocated in different parts of memory and are referenced through different variables.

Each piece of a file being arranged is defined as a “chunk.” We will assume that chunks are always the same size, meaning there is never a “leftover” chunk. In particular, the file size must be a multiple of the chunk size. The size of each chunk is specified by the user in stdin as a command line argument. The name of the output file is “output.txt”. The number of chunks that make up a file can be calculated by dividing the file size by the specified chunk size.

The order in which a file’s chunks are rearranged is based on a *rearrangement pattern*, similar to pattern matching. Each chunk is named in order with respect to the alphabet. A file with 5 chunks would have chunks ‘a’ to ‘e’. A possible reordering of this file would be the rearrangement pattern ‘a c b e d’. For example, given the file “AAABBBCCCDDDEEE”, a specified chunk size of 3, and the pattern ‘a c b e d’ the output of the program should be “AAACCCBBBEEEDDD”.




Figure 1: Assignment architecture

The rearrangement pattern is input via stdin as illustrated in the figure below. Here 4 indicates the chunk size and `numbers.txt` is the file to rearrange.

```

-> FileWriteAssignment java TextSwap 4 numbers.txt
Input 5 character(s) ('a' - 'e') for the pattern.
b
a
c
e
d
  
```

## 2.1 File Structure

This assignment consists of the following files:

- `Interval.java`. Declares the `Interval` class.
- `TextSwap.java`. Declares the `TextSwap` class. This class holds the main method from which the assignment is executed.
- `Swapper.java`. Declares the `Swapper` class.
- `letters.txt`. A sample text file.

### 2.1.1 TextSwap

The `TextSwap` class contains much of the logic for the assignment. It has the methods:

- `private static String readFile (String filename) throws Exception`

This method should read from a specified file by placing it into a `StringBuilder` object and then returning the `toString()` of that object.

- `private static Interval[] getIntervals (int numChunks, int chunkSize)`

This method returns an array of “Intervals”. An interval is just a pair of integers that specify the start and end index of a chunk. These intervals will be delegated to the appropriate threads to ensure the reordering is proper.

- `private static char[] runSwapper (String content, int chunkSize, int numChunks)`

This method does much of the actual logic of the class. It creates the intervals, runs the Swapper threads, and returns the reordered buffer that will be written to the new file.

- `private static void writeToFile (String contents, int chunkSize, int numChunks) throws Exception`

This method writes the buffer to a new file.

- `public static void main (String [] args)`

The main should parse two command line inputs, a chunk size and a filename. The size of the file and the number of chunks should then be calculated, and the new pattern of letters should be read from stdin. If the number of chunks is more than 26, then execution should halt with an error message “Chunk size too small”. If the file size is not a multiple of the chunk size or the chunk size is not positive, then execution should halt with an error message “Chunk size must be positive and file size must be a multiple of the chunk size”. Note that there may be other methods necessary to complete this class, but they can also be inlined into these methods.

### 2.1.2 Swapper

This class, which should implement the `Runnable` interface, will write to the buffer given its `Interval`. It has the fields `offset`, `interval`, `content`, `buffer`:

```

1 public class Swapper implements Runnable {
2     private int offset;
3     private Interval interval;
4     private String content;
5     private char[] buffer;
6     ...
7 }
```

Offset: specifies the starting index in the buffer where the content will be placed. Interval: specifies the starting and ending index of the content in the original file that is being swapped. Content: the entire original file in a `String` (called “Input Buffer” in Figure 1). Buffer: The shared `char[]` that the result is being written to.

- `public void run ()`

Write the specified content into the buffer. Helper methods may be used to retrieve the content and insert it into the proper spot in the buffer.

### 2.1.3 Interval

This is exactly what you would expect from any “Pair” class. There is even a `Pair` class in the Java library that can be used instead.

## 3 Your Task

Implement `Swapper.java` and the methods `runSwapper` and `getIntervals` in file `TextSwap.java`.

## 4 Submission Instructions

Submit a zip file named `hw1.zip` through Canvas containing all the files included in the stub but where all required operations have been implemented. It should unzip to files, not a folder. One submission per group. Place the name of the other team member as a canvas comment.